Compiling Facts into Applications

Samuel Lippert
Drivly Inc

Abstract

We define AREST, a compositional architecture unifying Backus's FFP/AST [1] with Fielding's REST [4] via Halpin's ORM 2 [5]. Domain knowledge is expressed as FFP objects in FORML 2 [6]; the representation function ρ maps each object to its executable; the system state is Backus's AST cell sequence; the persistent data is Codd's named set [3]. We prove seven properties: specification equivalence, completeness of state transfer, HATEOAS as projection, derivability, query adequacy, derivation completeness, and scope.

1 Definitions

Definition 1 (AREST)**.** An *Applicative REpresentational State Transfer* system is a tuple (O, ρ, D, P, sub) where:

- O is the set of FFP objects [1, Sec. 13.2]
- $\rho : O \to (O \to O)$ is the representation function [1, Sec. 13.3.2]
- D is the AST state, a sequence of cells [1, Sec. 14.3]
- $P = \uparrow FILE : D$ is the population, a named set of relations [3, Sec. 1.6] whose elements are ORM 2 elementary facts [5, Ch. 2]
- sub is the subsystem function dispatching inputs to state transitions, producing output as REST representations [4] with HATEOAS links

The primitive atoms of O include Backus's structural functions [1, Sec. 11.2.3]. Sequences represent functional forms via metacomposition [1, Sec. 13.3.2]:

$$(\rho\langle x_1, \ldots, x_n\rangle) : y = (\rho\, x_1) : \langle\langle x_1, \ldots, x_n\rangle, y\rangle \tag{1}$$

2 Operations

2.1 Domain Primitives

Each combining form [1, Sec. 11.2.4] is an FFP object with a domain interpretation:

FFP Object	Form	Domain
$\langle CONS, s_1, \ldots, s_n\rangle$	Construction	Graph Schema
s_i	Selector	Role
$\langle COMP, f, g\rangle$	Composition	Derivation
$\langle COND, p, f, g\rangle$	Condition	Constraint
$\langle \alpha, f\rangle$	Apply-to-all	Traversal
$\langle /, f\rangle$	Insert	Aggregation
$\langle bu, f, x\rangle$	Binary-to-unary	Query
$\bar{x}$	Constant	Literal

A graph schema $\langle CONS, s_1, \ldots, s_n\rangle$ applied to all roles is a fact; partially applied via bu it is a query.

2.2 Relational Algebra

Codd's adequate θ_1 [3, Sec. 2.2] as FFP definitions:

$$\pi_L(R) \equiv \alpha\,[s_{i_1}, \ldots, s_{i_k}] : R \tag{2}$$

$$R * S \equiv Filter(eq \circ [s_{sh}]) \circ distl \tag{3}$$

$$R_{L|_M} S \equiv Filter(p) : R \tag{4}$$

$$R \cdot S \equiv \pi_{1s}(R * S) \tag{5}$$

$$\gamma(R) \equiv Filter(eq \circ [s_1, s_n]) : R \tag{6}$$

2.3 State and Cells

The state D is a sequence of cells $\langle CELL, n, c\rangle$ [1, Sec. 14.3]:

$$\uparrow n : D \to c \qquad \text{(fetch)} \tag{7}$$

$$\downarrow n : \langle x, D\rangle \to D' \qquad \text{(store)} \tag{8}$$

$FILE$ is a cell whose contents is P. Each entity is a cell. Cells may contain sub-stores [1, Sec. 14.7]; partitioning is a cell containing a subset of P.

2.4 Subsystem Dispatch

Following [1, Sec. 14.4.2]:

$$\begin{aligned}
&sub \equiv \\
&\quad is\text{-}chg \circ \uparrow K \to [rpt, apl] \circ [\uparrow I, defs]; \\
&\quad is\text{-}qry \circ \uparrow K \to [qry \circ [\uparrow I, \uparrow F], defs]; \\
&\quad is\text{-}cmd \circ \uparrow K \to chk \circ apl \circ [\uparrow I, defs]; \\
&\quad is\text{-}upd \circ \uparrow K \to \\
&\qquad [rpt, \downarrow F \circ [upd, defs]] \circ [\uparrow I, \uparrow F]; \\
&\quad [err \circ [\uparrow K, \uparrow I], defs]
\end{aligned} \tag{9}$$

2.5 Commands

Each interaction is an AST transition [1, Sec. 14.3.1]:

$$\mu(SYSTEM : x) = \langle o, D' \rangle \tag{10}$$

A CREATE command:

$$create = emit \circ validate \circ derive \circ resolve \tag{11}$$

A state machine is a fold [1, Sec. 11.2.4]:

$$machine = /transition : \langle e_1, \ldots, e_n \rangle \tag{12}$$

$$transition = (p \to \bar{s}';\ id) \tag{13}$$

3 Execution Semantics

3.1 Evaluation as β-Reduction

FFP objects denote λ-terms via Backus's correspondence [1, Sec. 12.9]. The representation function ρ yields the executable term. Execution is applicative-order (innermost-first) β-reduction to normal form:

μ evaluates an expression under definitions D by repeatedly replacing its innermost application $(x : y)$ with $(\rho\, x) : y$, where ρ resolves x to a λ-term under environment D and primitives over P [1, Sec. 13.3.1]. Each step is a β-reduction: ρ yields the λ-term; application reduces it. The function $apply(func, x, D)$ in the implementation is this reduction step.

All evaluation in AREST is over finite P with monotonic derivations; reduction terminates for all requests.

3.2 Evaluation Strategy

Three constraints govern execution:

1. **Reduction order:** applicative (call-by-value). All arguments are reduced before the function is applied.
2. **Scope:** evaluation is confined to the request-local partition of D. RMAP [5, Ch. 17] determines the partition from the schema's uniqueness constraints.
3. **Derivation:** forward chaining, monotonic, evaluated to the least fixed point per request. Since *derive* only adds facts over finite P_{local}, the fixed point exists (Knaster-Tarski) and is reached in $\leq |P_{local,max}|$ iterations.

The completeness result (Theorem 3) assumes evaluation to the least fixed point under this strategy.

3.3 Boundary Conditions

The system operates over three categories:

1. **Facts** (P): ORM 2 elementary facts, persisted in cells.
2. **Computations** (O via ρ): total, deterministic functions over P.
3. **Inputs** (I): request payload, timestamp, external signals — part of the operand x in $\mu(SYSTEM : x)$ [1, Sec. 14.3.1].

Time, randomness, and external effects are classified as follows:

- **Time**: an input fact. "Now" is part of x, not hidden state. Time-dependent queries are constraints on timestamp facts in P.
- **Effects**: the output o in $\langle o, D' \rangle$. The engine computes directives (send email, charge payment); the runtime executes them. Outcomes are facts in D'.
- **Non-determinism**: identity generation is a primitive function producing a fresh atom, or the client provides the ID in I.
- **Partial information**: the world assumption per noun [5, Ch. 7] governs incompleteness. CWA: not stated = false. OWA: not stated = unknown.
- **Identity**: the reference scheme [5, Ch. 5] determines identity before facts are asserted. *resolve* in (11) applies the selector.

AREST governs the derivation of representations from (P, I). Theorem 8 holds: every value derivable from (P, I) by a computable function is in scope; values outside are non-factual or non-computable.

4 Properties

Theorem 1 (Grammar Unambiguity)**.** Let N be a declared set of object type names and F a declared set of predicate readings, such that no element of N contains a formal item of the grammar as a substring. Then each FORML 2 sentence over (N, F) has exactly one parse.

Proof. A FORML 2 sentence has the structure:

$$[modal]\ \ [quant_1]\ \ n_1\ \ r\ \ quant_2\ \ n_2\ \ [qual]$$

where $modal \in$ {"It is obligatory that", "It is forbidden that", "It is permitted that", ε}; $quant_1 \in$ {"Each", "For each", "the same", ε}; $n_1, n_2 \in N$; $r \in F$; $quant_2 \in$ {"at most one", "some", "exactly one", "at least k and at most m", "more than one", "no", ε}; $qual \in$ {"who ...", "that ...", ε}.

(1) The modal prefixes are pairwise non-overlapping strings, so *modal* is determined by the sentence prefix.

(2) After *modal*, the quantifier $quant_1$ is determined: "Each" vs. "For each" vs. "the same" vs. a name from N are distinguishable because no $n \in N$ contains a formal item as a substring (by hypothesis).

(3) n_1 is determined by longest-first matching against N. Since names do not contain formal items, the boundary between n_1 and r is unambiguous.

(4) r is the maximal token sequence between n_1 and $quant_2$. The quantifier phrases in $quant_2$ are pairwise distinct strings: {"at most one", "some", "exactly one", "at least", "more than one", "no"}. None is a prefix of another in this position. Since no $n \in N$ contains these phrases, the boundary between r and $quant_2$ is unambiguous.

(5) n_2 is determined by longest-first matching against N after $quant_2$.

(6) *qual* is determined by "who" or "that" (personal vs. impersonal pronouns per [7, Sec. 1.5]).

Each structural position admits exactly one parse. The constraint type is determined by the pair $(quant_1, quant_2)$, which maps injectively to constraint kinds: (Each, at most one) $\to$ UC; (Each, some) $\to$ MC; (Each, exactly one) $\to$ UC+MC; (Each, at least $k\ldots$) $\to$ FC; (the same, more than one) $\to$ UC$^-$. No two pairs map to the same kind. □

Theorem 2 (Specification Equivalence)**.** Let $parse : R \to \Phi$ and $compile : \Phi \to O$ be injective, with R the set of FORML 2 readings, Φ the set of FOL formulas, and O the set of FFP objects. Then:

$$parse^{-1} \circ compile^{-1} \circ compile \circ parse = id_R \tag{14}$$

and $\rho(compile(parse(r)))$ is the executable for every $r \in R$.

Proof. *parse* is injective by Theorem 1: each sentence has one parse, so distinct sentences yield distinct formulas. *compile* is injective: the $(quant_1, quant_2)$ pair determines the constraint kind (Theorem 1), which determines the FFP object structure (Condition for UC, null-test for MC, count for FC). Both have left inverses; composing: $parse^{-1} \circ compile^{-1} \circ compile \circ parse = id_R$. Executability: ρ maps the compiled object to a function by [1, Sec. 13.4]. □

Theorem 3 (Completeness of State Transfer)**.** *create* (11) produces a complete $\langle o, D' \rangle$ containing entity, state machine instance, derived facts, violations, and HATEOAS links.

Proof. Let $P = \uparrow FILE : D$. (1) *resolve* inserts entity facts: $P' = P \cup E$. (2) *derive* is monotonic over finite P; by Knaster-Tarski, a unique least fixed point P'' exists. (3) *validate* evaluates all constraint objects in D against P'': $V = \{(\rho\, c) : P'' \mid c \in D_{constraints}\}$. (4) *emit* produces $\langle o, D[FILE \mapsto P''] \rangle$ if no alethic violation, else $D' = D$. The state does not change during evaluation [1, Sec. 14.1]. □

Theorem 4 (HATEOAS as Projection)**.** Let $T = \{(s_i, s_j, e_k)\} \subseteq P$ be transition facts and s the current status. Then:

$$links(s) = \pi_{event}(Filter(p) : T) \tag{15}$$

where $p(t) = (s_{from}(t) = s) \vee anc(s_{from}(t), s)$.

Proof. $T \subseteq P$. *Filter* is restriction (4); π is projection (2). Both are θ_1 operations on the named set. The representation renders P; $links(s)$ is derived from P by θ_1 operations and therefore appears in the representation without external computation. □

Theorem 5 (Derivability)**.** For every domain value v in an AREST representation:

$$\exists f \in O.\ v = (\rho\, f) : P \tag{16}$$

Proof. Domain values partition into: (1) entity attributes (facts in P; $v = (\rho\, s_i) : fact$), (2) derived facts ($v = (\rho\, r) : P$ for derivation rule r), (3) state machine status (derived fact), (4) links ($v = (\rho\, \langle COMP, \pi, Filter(p) \rangle) : T$ by Theorem 4), (5) violations ($v = (\rho\, c) : P$ for constraint c). Each case yields (16). □

Theorem 6 (Query Adequacy)**.** Every query on P is expressible using θ_1, each element of which is an FFP definition (2)–(6).

Proof. P is a named set of relations in normal form (ORM 2 elementary facts are 5NF by construction [5, Ch. 12]). θ_1 is adequate for the named set [3, Sec. 2.2]. Each element of θ_1 is an FFP definition by (2)–(6). □

Theorem 7 (Derivation Completeness)**.** Every θ_1-derivable relation has a FORML 2 reading. Forward chaining evaluates the corresponding FFP objects to the least fixed point.

Proof. Each θ_1 operation has a FORML 2 verbalization: projection verbalizes as role selection ("X has Y"); join as shared-role composition ("X who R_1 Y and R_2 Z"); restriction as constraint predicate ("X where Y exceeds Z"); tie as self-reference ("X is ancestor of X"). By Theorem 1, each verbalization has a unique parse, so each θ_1 composition has a unique reading. Forward chaining applies derivation rules (FFP objects) to P monotonically; by Knaster-Tarski the least fixed point exists and is reached in $|P_{\max}|$ steps. □

Theorem 8 (Scope)**.** The architecture covers all values that are both factual and computable. Values outside are either non-factual or non-computable.

Proof. Any factual proposition (a statement with a truth value) is expressible as a tuple in a relation [3, Sec. 1.3], hence a fact in P. Any computable function (an effective method) is expressible as a λ-term [2]; Backus shows the equivalence of FP combining forms to λ-calculus combinators [1, Sec. 12.9], and FFP extends FP with metacomposition. Therefore every factual-and-computable value is $(\rho\, f) : P$ for some $f \in O$, $P \subseteq D$. Non-factual values (no truth value) cannot be relations. Non-computable values (no effective method) cannot be computed by any system. □

Corollary 1 (Algebraic Optimization)**.** Programs composed from FFP combining forms admit algebraic transformation via the laws of [1, Sec. 12.2]. Compilation from FFP objects to optimized evaluation preserves ρ-semantics.

Proof. The 24 laws of the FP algebra [1, Sec. 12.2] are equations between functional forms. An FFP object representing a program can be transformed to an equivalent object by applying these laws. The ρ-image is preserved because the laws hold for all objects in the domain. Ahead-of-time compilation from objects to evaluation trees is one such transformation. □

5 Example

```
Order(.OrderId) is an entity type.
Customer(.Name) is an entity type.
Order was placed by Customer.
  Each Order was placed by exactly one
  Customer.
Status 'Draft' is initial in State Machine
  Definition 'Order'.
Transition 'place' is from Status 'Draft'.
  Transition 'place' is to Status 'Placed
      '.
```

The constraint reading is an FFP object; ρ yields:

$$(\rho\, reading) = \langle COND,\ (count \leq 1) \circ Filter(eq \circ [s_1, \bar{o}]),\ \bar{\phi},\ violation\rangle : P_{placedBy}$$

`POST /orders {"customer":"acme"}` dispatches as *is-cmd* in (9). The engine applies *create* (11):

$$P' = P \cup \{Order(\texttt{ord-1}),\ placedBy(\texttt{ord-1}, \texttt{acme}),\ SM(\texttt{ord-1}, \texttt{Draft})\}$$

Response (Theorem 4: links from P'):

```
{ "id": "ord-1", "customer": "acme",
  "status": "Draft",
  "_links": {
    "place": {
      "href": "/orders/ord-1/transition",
      "method": "POST" } } }
```

Following "place" applies (12). The new representation contains $\{ship\}$, not $\{place\}$.

6 Conclusion

AREST composes FFP/AST [1] with REST [4] via ORM 2 [5]. Readings are FFP objects; ρ maps them to executables; the population is Codd's named set [3] stored as Backus's FILE cell; transitions produce hypermedia. The seven theorems establish: specifications and executables are the same object (Theorem 2); state transfer is complete in one application (Theorem 3); HATEOAS is a θ_1 projection (Theorem 4); all domain values are derivable (Theorem 5); all queries are expressible (Theorem 6); all derivations are verbalizable (Theorem 7); and the scope is factual-and-computable, which is everything (Theorem 8).

Acknowledgments

Claude (Anthropic) was used as a programming assistant during implementation and as a drafting aid during manuscript preparation. All architectural decisions and the composition of FFP/AST with REST were made by the human author.

References

[1] J. Backus, "Can Programming Be Liberated from the von Neumann Style? A Functional Style and Its Algebra of Programs," *Commun. ACM*, 21(8):613–641, 1978.

[2] A. Church, "An Unsolvable Problem of Elementary Number Theory," *Amer. J. Math.*, 58(2):345–363, 1936.

[3] E.F. Codd, "A Relational Model of Data for Large Shared Data Banks," *Commun. ACM*, 13(6):377–387, 1970.

[4] R.T. Fielding, "Architectural Styles and the Design of Network-based Software Architectures," Ph.D. dissertation, Univ. of California, Irvine, 2000.

[5] T. Halpin, *Information Modeling and Relational Databases*, 2nd ed. Morgan Kaufmann, 2008.

[6] T. Halpin and M. Curland, "Automated Verbalization for ORM 2," in *OTM Workshops*, pp. 1181–1190, 2006.

[7] T. Halpin, M. Curland, and CS445 Class, "ORM 2 Constraint Verbalization: Part 1," Technical Report ORM2-02, Neumont University, 2006.