

Compiling Facts into Applications

Samuel Lippert
Drivly Inc

Abstract

A single set of FORML2 declarations can define a database schema, state machine workflows, and a REST API with hypermedia navigation, producing a complete application from domain facts alone. We formalize this as AREST, a Backus AST system [1] whose definitions are FORML2 readings [8] that are compiled to a named set [3]. REST endpoints are applications of a single system function. Ingesting new readings is self-modification. Five theorems establish that the grammar is unambiguous, that specifications and executables are the same object up to a computable canonical form, that state transfer is complete, that HATEOAS links are projections of the population, and that every value in the API is derived from facts. Together, they localize informality to the registered platform layer, an enumerable surface that is itself a queryable fact set.

1 Terminology

Three terms carry overloaded abbreviations. *AST* follows Backus’s usage (Applicative State Transition, 1978), not the compiler term. *ORM* follows Halpin’s usage (Object-Role Modeling, 2001), not the persistence pattern. *FP* follows Backus’s usage (Function-level Programming), not the broader tradition. *FFP* is Backus’s extension of FP with a representation function ρ that maps objects to the functions they denote [1, Sec. 13]. This paper assumes comfort with the cited formalisms. The motivating examples are included to identify the construction, not to replace the references.

2 Definition

Definition 1 (AREST). An *Applicative REpresentational State Transfer* system is an AST system [1, Sec. 14] whose state D satisfies:

- The *FILE* cell contains a population $P = \uparrow \text{FILE} : D$, a named set of relations [3, Sec. 1.6] whose elements are ORM2 elementary facts [6, Ch. 3]
- The *DEFS* cells contain FFP objects [1, Sec. 13.2]: readings compiled from FORML2 declarations [8], and functions registered by the runtime

The *schema* S is the set of compiled objects in *DEFS*: fact type constructors, constraint objects, derivation

rules, and state machine definitions. The representation function ρ [1, Sec. 13.3.2] maps an object to the function it represents. Given a fact type, ρ looks it up in *DEFS* and returns the functional form. The system function $SYSTEM : x = \langle o, D' \rangle$ [1, Sec. 14.3.1] is defined by the definitions in D . REST endpoints are applications of *SYSTEM* with distinct inputs. A reference implementation is available [11].

3 Example

An order domain is declared in FORML2:

```
Order(.OrderId) is an entity type.
Customer(.Name) is an entity type.
Order is placed by Customer
  Each Order is placed by exactly one
  Customer.
Customer ships Order.
State Machine Definition 'Order' is for
  Noun 'Order'.
Status 'In Cart' is initial in
  State Machine Definition 'Order'.
Transition 'place' is defined in
  State Machine Definition 'Order'.
Transition 'place' is from
  Status 'In Cart'.
Transition 'place' is to
  Status 'Placed'.
Transition 'place' is triggered by
  Fact Type 'Customer places Order'.
Transition 'ship' is defined in
  State Machine Definition 'Order'.
Transition 'ship' is from
  Status 'Placed'.
Transition 'ship' is to
  Status 'Shipped'.
Transition 'ship' is triggered by
  Fact Type 'Customer ships Order'.
```

These readings compile to a database table, a foreign key, a uniqueness constraint, and a three-state machine driven by two trigger fact types, which is everything needed to run an order pipeline. Consider the constraint “Each Order is placed by exactly one Customer.” The compiler recognizes this as a uniqueness restriction on the *placedBy* role and emits:

$$(\rho \text{ reading}) : P_{\text{placedBy}} = \text{Filter}(\text{count}(s_1) > 1) : P_{\text{placedBy}}$$

The result is the set of Orders with more than one Customer in *placedBy*. An empty set means the constraint holds.

The operation `POST /orders {"customer":"acme"}` applies $SYSTEM:x$ where the addressed entity is *orders* and the operation is *create*. The pipeline produces:

$$P' = P \cup \{Order(ord-1, placedBy(ord-1, acme), SM(ord-1, In\ Cart))\}$$

The response includes links derived from P' :

```
{ "id": "ord-1", "customer": "acme",
  "status": "In Cart",
  "_links": {
    "place": {
      "href": "/orders/ord-1/transition",
      "method": "POST" } } }
```

Following “place” advances the state machine. The new representation contains $\{ship\}$, not $\{place\}$, because the status is now Placed. If a second command assigns a different Customer to the same Order, the restriction produces a non-empty violation set and the command is rejected. The violation is returned as the canonical reading of the violated constraint.

The constraint, the state machine, the API response, and the HATEOAS links above are different views of the same data. The mechanism that makes this possible is metacomposition. A fact type [6, Ch. 3] is a constructor of roles [1, Sec. 11.2.4]. Populate the roles with instances and you get a fact: $\langle ft, o_1, \dots, o_n \rangle$. Metacomposition [1, Sec. 13.3.2] resolves a fact by looking up its fact type:

$$(\rho\langle x_1, \dots, x_n \rangle):y = (\rho x_1):(\langle x_1, \dots, x_n \rangle, y) \quad (1)$$

The fact is a function: $\lambda f. f(o_1, \dots, o_n)$, accepting an operation and applying it to its own objects. The constraint reading, the state machine transition, and the API response above are all ρ -applications over the same facts.

4 The System Function

Each input x addresses an entity and carries an operation. The system function routes the input to the addressed entity and passes the operation as the operand:

$$SYSTEM:x = (\rho(\uparrow entity(x):D))\uparrow op(x) \quad (2)$$

The entity is fetched from D . ρ resolves it to a functional form. The operation is applied as the operand. The entity handles the dispatch, not the system function. New operations are registered in *DEFS* without modifying any entity, and new entity types are added without modifying any operation.

Remark (Specialization of Backus’s subsystem). Backus’s subsystem [1, Sec. 14.4.2] dispatches on a *KEY* cell with five clauses. AREST collapses them into a single ρ -dispatch (2): the entity determines the functional form,

and the HTTP method determines the operation class. Self-modification is the same mechanism: the addressed entity is *DEFS* and the operation is *compile* $\circ$ *parse*. The new FFP objects are stored via $\downarrow DEFS$, and subsequent applications of *SYSTEM* evaluate the new definitions.

4.1 Commands

μ denotes evaluation [1, Sec. 14.3.1]. Each application is an AST transition: the state before, the state after, nothing else changes.

A CREATE command:

$$create = emit \circ validate \circ derive \circ resolve \quad (3)$$

resolve determines identity from the reference scheme [6, Ch. 5]. *derive* forward-chains derivation rules to the least fixed point. *validate* applies every constraint as a restriction over the population, producing violation sets. *emit* always runs. An alethic violation rejects ($D' = D$). A deontic violation warns but succeeds ($D' = D''$). The violation message is the canonical reading (Corollary 1).

A state machine processes events chronologically:

$$machine(s_0, E) = foldl\ transition\ s_0\ (order_\tau E) \quad (4)$$

$$transition(s, e) = (p \rightarrow \bar{s}'; \bar{s}):\langle s, e \rangle \quad (5)$$

The transition function is a chain of conditions [1, Sec. 11.2.4]: p tests whether $\langle s, e \rangle$ matches a declared transition, f returns the target state, g tries the next or returns $\bar{s}$ as fallback. $order_\tau$ sorts E by each event’s occurrence timestamp τ , present for every recorded fact, since a fact *is* an event (ties, which intra-request cascades produce, break by chain position, the recorded derivation order, so the sort key is total and the fold stays deterministic for Corollary 5). Equation (4) is *reconstruction*: a replay of the whole stream from the initial state, for migration and audit. The *live* status is not recomputed on each read but maintained incrementally; each new event folds forward from the stored status (the cell fold, 20). Both are the one fold, differing only in base: s_0 over the ordered stream, or the stored state over the new events. Neither has side effects.

Facts as events. Creating a fact fires an event: a recorded fact *is* an event and a fact type *is* an event type (each a subtype), so the event carries the fact’s occurrence timestamp and its arguments are the role bindings. A transition that declares its trigger as a fact type fires automatically when that fact enters P . The metamodel reifies this as *Event caused Transition in State Machine*, joining each event to the transitions its event type triggers; the machine’s *currently-in Status* is the latest-wins fold of those transitions per machine, maintained incrementally. Reconstruction (4) is that fold replayed from s_0 ; the derivations expose the fold’s inputs and output as ordinary facts subject to constraints, derivation rules, and HATEOAS projection. During *create* (3), *resolve* and *derive* push facts into P ; each produces an event that the

machine fold (4) consumes. An entity whose facts satisfy all outgoing transition guards advances from its initial status to its terminal status in a single *SYSTEM* call. External events from a message queue, a webhook, or a peer enter the same stream without requiring a local fact: an inbound webhook is itself a fact (*Webhook Event Type yields Fact Type with Role from JSON Path*), and the runtime materializes the yielded facts into P before the fold runs. The machine does not distinguish between fact-fired and externally fired events; both are elements of E in the fold (4).

Deletion is a transition to a terminal state. The entity reaches a state with no outgoing transitions and is excluded from query results by restriction.

4.2 Platform Binding

The runtime registers functions into *DEFS* via $\downarrow DEFS$. A browser registers rendering functions; a server registers *httpFetch*, *upsert*, *notify*:

$$\begin{aligned}(\rho \text{ fact}): \text{render} &\rightarrow \text{widget} \\(\rho \text{ fact}): \text{httpFetch} &\rightarrow \text{response} \\(\rho \text{ fact}): \text{upsert} &\rightarrow D'\end{aligned}$$

Runtime functions that interact with external state execute during *resolve* (3), populating cells before *derive* runs. Functions that consume the representation execute after *emit*. This keeps *derive* pure and terminating over P . A single *SYSTEM* serves browser, server, and storage runtimes by varying *DEFS*, not by varying the logic.

The binding is live: when a cell’s contents change via $\downarrow$, every ρ -application over that cell re-evaluates. This is streaming. There is no need for a separate pub/sub layer, event bus, or callback registry. A subscriber is a ρ -application that has not yet been evaluated against the current D . When D changes, it evaluates. A sensor reading, a Kafka message, a derived inference, and a user click are all facts in P or events in E . They enter D through the same $\downarrow$ operator, and every bound ρ -application sees the update through the same mechanism.

4.3 Data Federation

A federated noun is one whose population is resolved by ρ at evaluation time from an external source rather than read from a local cell. The reading “Noun is backed by External System” compiles to a populating function *populate_n* registered in *DEFS*. During *resolve* (3), $\rho(\text{populate}_n)$ fetches from the external source and produces facts:

$$\rho(\text{populate}_n): I \rightarrow \{f_1, \dots, f_k\} \subseteq P_{OWA}$$

The fetched facts enter P under OWA (absence is unknown, not false). Constraints and derivation rules evaluate over the unified P without distinguishing local from

federated facts. A subset constraint can span a local entity and a federated one. No schema translation is required because the fact type declared in the readings is the response schema, and JSON keys map to role bindings through the compiled schema.

The compile step constructs *populate_n* from readings: the external system’s URL, the noun’s URI path, authentication headers, and credential references. SSRF validation rejects internal or loopback URLs at compile time, before any fetch can execute [11]. Cell isolation (Definition 2) ensures federated facts are cached for the request duration so that two concurrent μ applications do not re-fetch.

The approach requires no ETL, no integration middleware, and no separate query language for external data. The external system is a ρ -application, and the rest of the algebra does not change.

4.4 State and Cells

The state D is a sequence of cells $\langle \text{CELL}, n, c \rangle$ with Backus’s fetch ($\uparrow n: D \rightarrow c$) and store ($\downarrow n: \langle x, D \rangle \rightarrow D'$) operators [1, Secs. 13.3.4, 14.3]. *FILE* is a cell whose contents is P . RMAP [6, Ch. 10] determines which facts belong to which cell from the schema’s uniqueness constraints: the result is a 3NF row, the complete set of facts that depend on one entity’s key. Each entity is a cell. Cells may contain sub-stores [1, Sec. 14.7].

Definition 2 (Cell Isolation). For each cell $\langle \text{CELL}, n, c \rangle$ in D , at most one μ application that writes to n may be in progress at any time. Concurrent μ applications over disjoint cells are permitted.

RMAP assigns each entity to its own cell whose fold is independent. A single-threaded runtime satisfies Definition 2 trivially. A distributed system requires one writer per cell. Without isolation, two concurrent *create* operations on the same cell may both pass *validate* against stale P and produce inconsistent state without an error.

Definition 3 (Constraint Scope). The *scope* of a constraint c is the RMAP image of the fact types its restriction ranges over: $\text{scope}(c) = \text{RMAP}(\text{factTypes}(c))$, the set of cells those fact types can populate. Two commits *conflict* on c when both write cells in $\text{scope}(c)$.

The scope is schema-level, computable at compile time, and deliberately conservative: it contains every cell the restriction could ever read, not the cells a particular evaluation happened to read. An instance-relative read set would under-capture phantoms; the cell for an Order that does not exist yet is read by nobody, and it is exactly the insert that constitutes the skew. Cell Isolation serializes only the constraints whose scope is a single cell; the constraints worth having span cells, and there the model admits write skew. For “each Customer places at most 3 Orders,” two concurrent creates in distinct Order cells each validate against a P that lacks the other’s

pending insert, and both commit, jointly violating a constraint neither violated alone. The model therefore imposes an obligation rather than resolving one: commits that conflict on a constraint's scope must serialize. The cardinality family discharges it inside Definition 2 itself: the count groups by the constrained key (the Customer, above), so the serialization point is that entity's existing cell, single-writer-per-scope-key as a generalization of Cell Isolation rather than a new mechanism. The general cases are commit-time conflict detection over the scope, or the total event order of Corollary 5. A single-threaded runtime discharges all of it trivially. Theorem 3 assumes scope-serialized commits below.

5 Formal Foundations

Section 3's metacomposition identity (1) is the domain interpretation of one FFP primitive. This section enumerates the correspondence that made the four views of Section 3 necessarily the same object.

5.1 Domain Primitives

Each combining form [1, Sec. 11.2.4] is an FFP object with a domain interpretation:

FFP Object	Form	Domain
$\langle CONS, s_1, \dots, s_n \rangle$	Construction	Fact Type
s_i	Selector	Role
$\langle COMP, f, g \rangle$	Composition	Derivation
$\langle COND, p, f, g \rangle$	Condition	Guard
$\langle \alpha, f \rangle$	Apply-to-all	Traversal
$\langle /, f \rangle$	Insert	Aggregation
$\langle bu, f, x \rangle$	Binary-to-unary	Query
$\bar{x}$	Constant	Literal
$distl$	Distribute-left	Join preparation
eq	Equality test	Predicate
$compact$	Remove $\perp$ elements	Set formation

Table 1: FFP primitives with domain interpretations.

A fact type $\langle CONS, s_1, \dots, s_n \rangle$ is a function that produces a fact when applied to instances. Partially applied via bu it produces a query.

Two derived forms:

$$Filter(p) \equiv compact \circ \alpha(p \rightarrow id; \perp) \quad (6)$$

α applies the condition to each element, yielding $\perp$ where p does not hold. $compact$ removes $\perp$ elements, producing a proper subsequence. FFP systems choose their primitive set [1, Sec. 13.4]; $compact$ is an AREST primitive analogous to APL compress.

$$\begin{aligned} foldl f z \langle \rangle &= z \\ foldl f z \langle e, E' \rangle &= foldl f (f z e) E' \end{aligned} \quad (7)$$

$foldl$ is derivable from Backus's *while* [1, Sec. 11.2.4].

The compiled form for each FORML 2 construct:

FORML 2 Construct	Compiled FFP Object
Entity type $X(.ref)$	$\langle CONS, s_1, \dots, s_n \rangle$
Fact type $X r Y$	$\langle CONS, s_1, s_2 \rangle$ (binary)
Derivation rule	$\langle COMP, f, g \rangle$
State machine	$foldl transition s_0 E (7)$
Instance fact	$\bar{x}$ asserted into P
<i>Constraints (all compile to $Filter(p) : P (6)$)</i>	
<i>Cardinality family: $p = count(scope) \notin [k, m]$</i>	
UC, MC, FC, XC, XO, OR	$scope \in \{role, participation\}$
<i>Membership family: $p = match(target, polarity)$</i>	
IR, AS, SY, AT (self-ref)	target = same population, reversed
IT, TR, AC (path)	target = transitive closure step
SS, EQ (set comparison)	target = cross-fact-type tuples
VC (value)	target = declared value set

Every ORM2 constraint compiles to a restriction (10) whose predicate belongs to one of two families. The *cardinality* family checks whether a count falls within bounds $[k, m]$. The *membership* family checks whether a tuple exists (or is absent) in a target set. Polarity is the only parameter distinguishing symmetric from asymmetric, subset from exclusion.

5.2 Relational Algebra

Codd's adequate θ_1 [3, Sec. 2.2] as FFP definitions:

$$\pi_L(R) \equiv \alpha[s_{i_1}, \dots, s_{i_k}] : R \quad (8)$$

$$R * S \equiv Filter(eq \circ [s_{sh}]) \circ distl \quad (9)$$

$$R_{L|M} S \equiv Filter(p) : R \quad (10)$$

$$R \cdot S \equiv \pi_{1s}(R * S) \quad (11)$$

$$\gamma(R) \equiv Filter(eq \circ [s_1, s_n]) : R \quad (12)$$

where s_{sh} selects the shared roles. P is a named set of relations in normal form (ORM2 elementary facts are 5NF by construction [6, Ch. 12]). Codd's claim, with its own scoping intact: "For the named set of relationships in non-inferential systems, it appears that an adequate collection θ_1 contains the following operations: projection, natural join, tie, and restriction" [3, Sec. 2.2]. Two qualifiers carry weight. "It appears that": adequacy is Codd's judgment, never proven, and we use it as such. "Noninferential": θ_1 is scoped to retrieval, and that is exactly where AREST uses it. The navigation layer is retrieval over the elementary-fact populations as the named relations, and it stays inside Codd's restriction (every predicate in Theorem 4 composes eq with selectors and constants). The inferential layer lives in composition: derivation is $\langle COMP, f, g \rangle$, precisely on the other side of the boundary Codd's carve-out drew for θ_1 . The constraint layer generalizes the restriction predicate (the uniqueness example of Section 3 filters on $count(s_1) > 1$), and counting sits outside the algebra Codd judged adequate; the algebra-with-aggregates equivalence is Klug's [9]. Navigation is θ_1 in Codd's own operators; constraints are θ_1

with an FFP predicate parameter; derivation is composition. The combining forms that query the population are the same combining forms that evaluate constraints and derive facts. There is no separate query language.

5.3 Execution Semantics

μ evaluates by applicative-order β -reduction [1, Sec. 13.3.1], repeatedly replacing $(x : y)$ with $(\rho x) : y$ under definitions D . Three constraints govern execution:

1. **Reduction order:** applicative (call-by-value).
2. **Scope:** confined to the request-local partition of D determined by RMAP.
3. **Derivation:** forward chaining, monotonic (rules are positive by the grammar’s derivation family; negation is confined to constraint rendering), evaluated to the least fixed point per request. Derivation only adds facts over finite P ; the fixed point exists by Knaster-Tarski and is reached in $\leq |P_{\max}| - |P|$ steps. The derivation chain is recorded and available on demand.

Remark (Cost model). The fixed point is the specification, not the execution plan. Because P is append-only (deletion is a terminal state, Corollary 3), $lfp(FS, \cdot)$ admits maintenance as a materialized view across the event stream rather than recomputation per request: insertion-only semi-naïve evaluation fires a rule only on the previous round’s delta [2], and the one controlled non-monotonicity, the current-status projection a new event logically retracts, is classical incremental view maintenance [5]. Per-request cost is then proportional to the new derivations and the representation touched, not to a re-chase of the partition (which RMAP already confines); Knaster-Tarski guarantees the cheaper path agrees with the specification.

The system operates over facts (P), computations (O via ρ), and inputs (I). Time is an input fact, not hidden state. Effects are directives in the output. Identity follows from the reference scheme; *resolve* (3) applies it. The world assumption per noun [6, Ch. 7] governs incompleteness: under CWA, a fact not in P is false; under OWA, it is unknown.

Remark (World assumption on populating functions). All populating functions are runtime functions registered in *DEFS*. A function that exhaustively enumerates its domain operates under CWA. One that cannot operate under OWA. Deontic constraints over OWA-populated facts are sound but not complete: a reported violation is genuine, but the absence of a violation is not a guarantee. This is the population-side view of the registered boundary that Corollary 2 makes enumerable.

6 Properties

The five theorems form a chain. Each depends on the previous one and eliminates one source of failure in the

path from domain knowledge to running software. They are not all of one register, and we say so up front: Theorems 1–4 have proof content, while Theorem 5 is true by construction (a closure invariant of the definition of *repr*) and is stated in the chain because it is the guarantee the agent argument consumes. The genre is Fielding’s: an architecture justified by exhibiting its own construction.

Theorem 1 (Grammar Unambiguity). Let N be a declared set of object type names and F a declared set of predicate readings, such that no element of N contains a formal item of the grammar as a substring. Then each sentence of the AREST fragment of FORML 2 over (N, F) (the three structural families below, which are exactly the sentences the compiler accepts) has exactly one parse.

Proof. FORML 2 sentences fall into three structural families, each determined by its prefix:

(a) *Quantified constraints.* Structure:

$[modal] \text{ quant}_1 n_1 r \text{ quant}_2 n_2 [qual]$
 where $modal \in \{\text{“It is obligatory that”, “It is forbidden that”, “It is permitted that”, “It is impossible that”, “It is possible that”, } \varepsilon\}$; $\text{quant}_1 \in \{\text{“Each”, “For each”, “the same”, “No”, } \varepsilon\}$; $\text{quant}_2 \in \{\text{“at most one”, “some”, “exactly one”, “at least } k \text{ and at most } m\text{”, “more than one”, “no”, } \varepsilon\}$. Noun positions are resolved by longest-first matching. Each structural position admits exactly one parse because (i) modal prefixes are pairwise non-overlapping, (ii) quantifier phrases are pairwise distinct, (iii) no $n \in N$ contains a formal item as a substring, and (iv) longest-first matching is deterministic. The constraint kind is determined by $(\text{quant}_1, \text{quant}_2)$.

(b) *Conditional constraints.* The “If . . . then” frame splits unambiguously on “then”. Ring constraints are identified by all noun bases being the same type. Subset constraints by existential “some” in the antecedent and anaphoric “that” in the consequent. Derivation rules occupy the remainder, and “remainder” is a partition claim, not a default: the two constraint markers are pairwise-disjoint syntactic criteria, a conditional sentence carrying neither marker parses as a derivation rule only if its consequent matches a declared predicate reading, and a conditional matching none of the three is rejected by the compiler rather than guessed at. The accepted conditional sentences are partitioned by construction.

(c) *Multi-clause constraints.* Keyword-delimited: “if and only if” (EQ), “exactly one of the following holds” (XO), “at most one of the following holds” (XC), “at least one of the following holds” (OR). The keyword is unambiguous.

The three families are distinguished by their sentence prefix at position (1). Within each family, the parse is deterministic. The reference implementation [10] constructively demonstrates unambiguous round-trip parsing for all 17 constraint kinds. $\square$

The hypothesis on N is enforced by the compiler, which rejects any declared name containing a formal item of

the grammar as a substring. Positivity of the derivation family is enforced the same way: a conditional whose antecedent carries negation (the “No” quantifier, “no” in a clause) is a constraint or a rejection, never a derivation rule, so the rules reaching F_S (Theorem 3) are positive by construction rather than by convention. The fragment is a proper subset of Halpin’s full FORML2: mixfix predicates of arbitrary arity, pronoun-correlated clauses, and nested objectification lie outside it and are rejected rather than guessed at. Ambiguity in controlled natural languages lives in attachment and scope, not tokenization; the families above are chosen so that neither arises, and the theorem claims nothing about sentences outside the fragment.

Theorem 2 (Specification Equivalence). Let $parse : R \rightarrow \Phi$ and $compile : \Phi \rightarrow O$ be the parse and compile maps, with R the set of readings in the AREST fragment of FORML2 over (N, F) (Theorem 1), Φ the set of FOL formulas, and O the set of FFP objects. Define *reading equivalence* as the kernel of the pipeline: $r_1 \sim r_2$ iff $compile(parse(r_1)) = compile(parse(r_2))$. Then $compile \circ parse$ is injective on $R/\sim$, and the quotient has a computable section that lands inside the grammar: with *verbalize* the Halpin–Curland primary-reading verbalization [7],

$$nf = verbalize \circ compile \circ parse \quad (13)$$

is an idempotent normalizer ($nf \circ nf = nf$) with $nf(r) \sim r$ for every $r \in R$, and $\rho(compile(parse(r)))$ is the executable for every $r \in R$.

Proof. Injectivity on the quotient is immediate: $R/\sim$ is the kernel quotient of $compile \circ parse$. The content is the section. *parse* is total and well-defined on R by Theorem 1: one parse *per sentence*, not one sentence per formula. Inverse readings such as “Order is placed by Customer” / “Customer places Order” are distinct elements of R with the same compiled object, which is exactly what $\sim$ absorbs. *compile* maps each parse to a restriction (10) whose predicate comes from one of two families parameterized by the constraint tuple. *verbalize* emits the primary reading of a compiled object [7]; the primary reading is a sentence of the fragment (witnessed for all 17 constraint kinds by the round-trip [10]), so it parses by Theorem 1, and it re-compiles to the object it verbalizes. Hence $nf(r) \sim r$, nf fixes exactly the canonical readings, and $nf \circ nf = nf$. Executability: ρ maps the compiled object to a function by [1, Sec. 13.4]. $\square$

The executable is the equivalence class; the primary reading is its canonical name; surface variants are presentations of the same object. “Specification and executable are the same object” is exact in this sense.

Remark (Surjectivity). *compile* is not surjective. Runtime-registered functions occupy the complement. FORML2 readings cover the domain layer; registered

functions cover the platform layer. Together they span *DEFS*, and the partition itself is a fact (Corollary 2).

Theorem 3 (Completeness of State Transfer). Let $F_S : 2^P \rightarrow 2^P$ be the monotonic operator that applies all derivation rules in schema S to a population (positive by the grammar: the derivation-rule family of Theorem 1 compiles positive bodies, and negation is confined to constraint rendering), and let C_S be the set of constraint objects in S . Under Definition 2 with scope-serialized commits (Definition 3), *create* (3) applied to input I and state D produces $\langle o, D' \rangle$ such that:

1. $P' = P \cup resolve(S, I)$ contains all entity facts entailed by (S, I) ;
2. $P'' = lfp(F_S, P')$, the least fixed point of F_S above P' ;
3. $V = \bigcup \{(\rho c) : P'' \mid c \in C_S\}$ collects all violations;
4. o is the REST representation of $(P'', V, links(s))$ and $D' = D[FILE \mapsto P'']$ if V contains no alethic violation, else $D' = D$.

Proof. P is finite. F_S is monotonic because every rule is positive; a rule with negation under a CWA noun would break monotonicity, and the grammar’s derivation family excludes it (stratified or well-founded extensions are future work, outside this theorem). By Knaster-Tarski, $lfp(F_S, P')$ exists and is unique. *validate* applies every $c \in C_S$ as a restriction over the fixed point. *emit* constructs o from P'' , V , and $links(s)$ (Theorem 4). Definition 2 ensures no concurrent write to the same cell, and scope serialization (Definition 3) ensures no write skew across the cells a constraint reads. $\square$

Theorem 4 (HATEOAS as Projection). All links in the representation are θ_1 operations on P and S .

(a) *Transition links.* Let $T = \{(s_i, s_j, e_k)\} \subseteq P$ be transition facts and s the current status. Then:

$$links(s) = \pi_{event}(Filter(p) : T) \quad (14)$$

where $p(t) = s_{from}(t) \in \{s\} \cup supertypes(s)$. Transition graphs may contain cycles. It is obligatory that each cycle has some exit transition, a deontic constraint that ensures liveness without requiring structural acyclicity.

(b) *Navigation links.* Let F be the set of all fact types in S and let U_f denote the role set spanned by the UC on fact type f . Let $S_f = R_f \setminus U_f$ be the non-spanned roles. For a noun n participating in f :

1. **Related collection.** f is a related collection on n , filtered by n . Always applies.
2. **Parent/child.** If $S_f \neq \emptyset$, nouns in S_f are children of nouns in U_f .
3. **Peer resource.** If $S_f = \emptyset$ (spanning UC), f is a navigable peer resource.

For the binary case:

$$children(n) = \pi_{noun(r_c)}(Filter(uc(r_c) \wedge eq \circ [noun(\bar{r}_c), \bar{n}]) : F_B) \quad (15)$$

$$parent(n) = \pi_{noun(\bar{r}_c)}(Filter(uc(r_c) \wedge eq \circ [noun(r_c), \bar{n}]) : F_B) \quad (16)$$

$$nav(e, n) = children(n) \cup parent(n) \cup related(n) \cup peers(n) \quad (17)$$

The full link set:

$$links_{full}(e, n, s) = nav(e, n) \cup links(s) \quad (18)$$

Proof. (a) $T \subseteq P$. *Filter* is restriction (10); π is projection (8). Both are θ_1 . (b) Each component of *nav* is a projection over a restriction. The predicate composes *eq* with selector and constant, all primitives from Table 1. All components of *links_{full}* are θ_1 operations on S and P . $\square$

Remark (Prior formal models). Finite-state models of RESTful systems exist [12]; they model the client’s traversal as a machine over received representations. Theorem 4 works on the other side of the interface: it derives the affordance set itself from the schema’s uniqueness-constraint structure and the population, the part of hypermedia Fielding left informal.

Remark (Selection). The theorem enumerates affordances and deliberately says nothing about choosing among them: when $|links(s)| > 1$, an agent needs a choice function, and the formalism leaves that joint open by design. Any principled selector that reads facts (weights, preferences, valuation fact types over the link set) sits in the constraint register of Section 5.2: a restriction with an FFP predicate parameter, since ranking is ordering and aggregation, not a dyadic comparison. Selection extends the architecture from inside that register rather than escaping the formalism.

Theorem 5 (Derivability). Define the *representation* of an entity e in population P under schema S as:

$$\begin{aligned} repr(e, P, S) = \{ & (\rho s_i) : f \mid f \in P, f \text{ involves } e \} \\ & \cup \{ (\rho r) : P \mid r \in S_{derivations} \} \\ & \cup \{ (\rho c) : P \mid c \in S_{constraints} \} \\ & \cup links_{full}(e, n, status(e, P)) \end{aligned}$$

Then for every $v \in repr(e, P, S)$, there exists $f \in O$ such that $v = (\rho f) : P$. Equivalently: no value in the REST representation is produced outside ρ .

Proof. Each component is a ρ -application by construction. Selectors extract attributes. Derivation rules compute derived facts. Constraints produce violation sets via restriction. Links are θ_1 operations (Theorem 4). No value in *repr* is produced outside ρ . $\square$

Corollary 1 (Violation Verbalization). If $(\rho c) : P'' \neq \emptyset$ for constraint c compiled from reading r , the violation output is *verbalize*(c) = *nf*(r) (Theorem 2). The error message is the canonical reading, deterministic regardless of which surface variant of r the author typed.

Halpin and Curland [7] give the automated verbalization algorithm. Composed θ_1 operations verbalize via FORML 2 qualified readings, which parse unambiguously by Theorem 1.

Corollary 2 (Enumerable Trust Boundary). Let every definition entering *DEFS* carry the fact type *Definition has Origin*, *Origin* $\in \{\text{compiled}, \text{registered}\}$, recorded by $\downarrow DEFS$ at entry. The informal surface of the system is

$$Filter(eq \circ [s_{origin}, \overline{\text{registered}}]) : DEFS,$$

a θ_1 restriction over *DEFS*-as-population, itself a ρ -application (Theorem 5).

The surjectivity remark (Theorem 2) and the world-assumption remark (Section 5.3) are the same boundary seen from the definition side and the population side: both name the compiled/registered partition of *DEFS*. Making the partition a fact type turns two concessions into an audit query. Informality is not eliminated; it is localized to an enumerable surface that is itself a queryable fact set, self-describing by construction. The same query names the kernel: *parse*, *compile*, and *verbalize* are themselves registered code that the theorems quantify over but do not verify, an LCF-style trusted kernel. The trusted base is exactly the compiler plus the registered layer, and the restriction above enumerates it.

Corollary 3 (Deletion as Terminal State). A state with no outgoing transitions produces $links(s) = \emptyset$ by Theorem 4. The entity is excluded from query results by *Filter*(*p_{live}*) : P .

Corollary 4 (Closure Under Self-Modification). After applying *SYSTEM* with operation *compile* $\circ$ *parse* to readings R' , Theorems 1–5 hold for all subsequent applications over D' .

Proof. Theorem 1 depends on the grammar, not on D . Theorem 2 depends on the grammar and the kernel quotient, neither of which varies with D . New definitions enter *DEFS* via $\downarrow DEFS$. Theorems 3–5 operate over P and S , both of which include the new content. No theorem’s proof depends on D being fixed at initialization. $\square$

Remark (Self-hosting). This is the property the formalism turns on: AREST ingests schemas that extend its own schema, and every theorem in the chain continues to apply. The system is not merely specified by readings; it is extended by them.

Remark (Migration). The compile op is itself a *SYSTEM* application and subject to *validate* (3). If an ingested schema introduces an alethic constraint that P

violates, the ingestion is rejected ($D' = D$). Migration readings that transform P must therefore land in the same ingestion (as derivation rules, transition triggers, or deontic rather than alethic constraints) or as a staged sequence whose intermediate readings produce compliant P before the stricter constraint is introduced.

Corollary 5 (Constraint Consensus). For two peers sharing population P , one proposes event e and the other applies *validate* (3) to $P \cup \text{resolve}(S, e)$. The event is accepted iff V contains no alethic violation. The constraint set C_S is the consensus predicate.

For $n > 2$ peers, a total ordering over the shared event log is additionally required. Safety follows from deterministic replay; liveness from the single-writer cell (Definition 2).

7 Distributed Evaluation

Each cell stores the 3NF row. The event stream is demultiplexed into per-cell streams by RMAP:

$$E_n = \text{Filter}(eq \circ [RMAP, \bar{n}]) : E \quad (19)$$

Each cell folds its stream independently:

$$D'_n = \text{foldl } \mu_n \ D_n \ E_n \quad (20)$$

where μ_n is *SYSTEM* restricted to cell n . Cross-cell queries read committed state:

$$P = \bigcup_n \uparrow \text{FILE} : D_n \quad (21)$$

Adding cells to the cluster scales *SYSTEM* horizontally without changing its definition.

A browser runtime is a node in the same model. Cells replicated locally fold events with zero latency, syncing asynchronously. Because the fold is deterministic, two nodes folding the same events reach the same state (Corollary 5). Local constraints (cardinality, value) evaluate immediately. Cross-cell constraints (subset, external uniqueness) evaluate against committed state after sync, and their commits serialize per Definition 3. Local evaluation is optimistic feedback, not a security boundary: acceptance is decided at commit, where *validate* runs against the authoritative P (Corollary 5).

For anonymous peers, events carry cryptographic signatures for identity, schema modifications require authorization (a deontic constraint), and every peer validates independently (Corollary 5). The hash chain provides ordering, constraints provide validity, signatures provide identity, and authorization provides permission, all composed from the existing formalism: validity is Theorem 3, representation is Theorem 5.

8 Middleware Elimination

Traditional middleware (authentication, validation, rate limiting, transformation) is either a restriction over P , a

composition, or a fact in P . Authorization is a derivation: “User accesses Domain iff User belongs to Organization that manages Domain” compiles to $\langle COMP, f, g \rangle$ whose ρ -image computes the access predicate over P . External service calls, database persistence, and UI rendering can be expressed as runtime functions registered into *DEFS* and applied via ρ . Theorem 5 leaves no need for a computation layer outside ρ .

9 Conclusion

The five theorems form a dependency chain from unambiguous grammar to full derivability, each eliminating one source of failure between domain knowledge and running application. Self-modification preserves all five (Corollary 4). Fielding’s remaining architectural constraints [4] follow: statelessness holds because each application of *SYSTEM* is a pure transition on D ; cacheability because the representation is a deterministic function of P and S (Theorem 5); layered system because cells are location-independent; code-on-demand because self-modification via ingesting readings preserves the theorems.

The five properties compose into a guarantee for automated agents. An LLM writing FORML2 readings produces unambiguous specifications (Theorem 1) that are the executables (Theorem 2). An LLM consuming the API reads values that are ρ -applications over P (Theorem 5) and discovers every valid action via HATEOAS links (Theorem 4). There are no undocumented endpoints, no hidden middleware, no informal contracts to hallucinate through. The forward chainer evaluates derivation rules to the least fixed point (Theorem 3), and the LLM reads conclusions, not premises. The logical chain was evaluated deterministically by the engine. When the chain itself is needed, the derivation trace is available on demand.

The algebraic laws of FP [1, Sec.12.2] are equations between functional forms. Since FFP objects represent the same forms via ρ , an FFP object can be transformed to an equivalent object by applying these laws. Ahead-of-time compilation from objects to optimized evaluation trees is one such transformation, and so is semi-naive evaluation [2]: the cost model of Section 5.3 is a program transformation in the same genre.

Common operational concerns map directly to the formalism: versioning is the event stream, migration is ingestion of new readings, backward compatibility is preserved by construction, and auditability is the derivation trace.

The gap between domain knowledge and running software is the gap where requirements drift, implementations diverge, and production breaks. AREST eliminates the gap by identifying what was never separate to begin with.

Acknowledgments

Claude (Anthropic) was used as a programming assistant during implementation and as a drafting aid during manuscript preparation. All architectural decisions and the composition of FFP/AST with REST were made by the human author.

References

- [1] J. Backus, “Can Programming Be Liberated from the von Neumann Style? A Functional Style and Its Algebra of Programs,” *Commun. ACM*, 21(8):613–641, 1978.
- [2] F. Bancilhon and R. Ramakrishnan, “An Amateur’s Introduction to Recursive Query Processing Strategies,” in *Proc. ACM SIGMOD*, pp. 16–52, 1986.
- [3] E.F. Codd, “A Relational Model of Data for Large Shared Data Banks,” *Commun. ACM*, 13(6):377–387, 1970.
- [4] R.T. Fielding, “Architectural Styles and the Design of Network-based Software Architectures,” Ph.D. dissertation, Univ. of California, Irvine, 2000.
- [5] A. Gupta, I.S. Mumick, and V.S. Subrahmanian, “Maintaining Views Incrementally,” in *Proc. ACM SIGMOD*, pp. 157–166, 1993.
- [6] T. Halpin, *Information Modeling and Relational Databases*, Morgan Kaufmann, 2001.
- [7] T. Halpin and M. Curland, “Automated Verbalization for ORM 2,” in *OTM Workshops*, pp. 1181–1190, 2006.
- [8] T. Halpin and J.P. Wijnenga, “FORML 2,” in *Enterprise, Business-Process and Information Systems Modeling (EMMSAD)*, LNBIP 50, pp. 247–260, 2010.
- [9] A. Klug, “Equivalence of Relational Algebra and Relational Calculus Query Languages Having Aggregate Functions,” *J. ACM*, 29(3):699–717, 1982.
- [10] T. Halpin and M. Curland, NORMA (Natural Object-Role Modeling Architect), open-source reference implementation, <https://github.com/ormsolutions/NORMA>.
- [11] S. Lippert, AREST, open-source reference implementation, <https://github.com/graphdl/arest>.
- [12] I. Zuřak, I. Budiselić, and G. Delač, “Formal Modeling of RESTful Systems Using Finite-State Machines,” in *Proc. ICWE*, LNCS 6757, pp. 346–360, 2011.